

# Distributional Split Criteria for Random Forests: Extensions, Shrinkage, and the Robustness of Mean Splitting

Silas Koemen

June 30, 2026

## Abstract

Distributional random forests replace mean-based CART splitting with criteria that compare the full conditional response distribution in candidate children. We implement and systematically study a family of such criteria inside a single honest-forest implementation: isotropic random-Fourier-feature maximum mean discrepancy (MMD), an anisotropic diagonal-bandwidth variant, an adaptive per-split frequency-selection variant, and a non-kernel sliced-Wasserstein criterion, together with post-hoc kernel-mean shrinkage of the forest weights. Using paired-seed comparisons across synthetic quantile mechanisms, real univariate benchmarks, a California-housing subsample curve, and multivariate synthetic and real responses, we characterize where each extension pays. Three findings recur. First, among distributional criteria ordinary isotropic MMD is already close to best in class: the anisotropic, adaptive-frequency, and sliced-Wasserstein extensions, and post-hoc shrinkage, do not systematically improve on it. Second, on scalar tabular regression mean-based CART splitting remains the robust default and wins many cells. Third, multivariate responses are the regime where distributional splitting clearly earns its keep, most sharply on a pure-dependence copula where the energy score separates the criteria even though marginal CRPS does not. The evidence supports a simple allocation story: distributional splitting helps only when non-location structure is both present and estimable; otherwise it dilutes split-selection power away from the mean. All criteria, the honest forest, and the paired-comparison harness are implemented in the open-source `drforest` library, whose Rust-backed split search makes broad criterion sweeps inexpensive.

## 1 Question

Random forests for distributional prediction estimate not just a conditional mean, but a conditional law. Quantile regression forests do this with standard CART tree structure and distributional leaf weights [2]. Distributional random forests instead propose splitting rules that compare the full response distributions in candidate children [1]. This is a natural generalization, especially for multivariate responses where a scalar ordering or separate marginal quantile fits are unsatisfactory.

This note studies that design space directly. Inside one honest-forest implementation we add four distributional split criteria—*isotropic MMD*, a diagonal-bandwidth anisotropic variant, an adaptive per-split frequency-selection variant, and sliced Wasserstein—and a family of post-hoc shrinkage variants, and compare them by paired seed. Two questions organize the study: *among distributional criteria, does any extension beat the plainest MMD?* and *against mean splitting, when is distributional splitting worth its variance?* The answers are, respectively, “not systematically” and “for multivariate responses and a few structured scalar problems,” with CART the robust default elsewhere. We test the multivariate regime explicitly, since it is what motivates distributional forests in the first place.

## 2 Split Criteria as Signal Allocation

For a node with observations  $(X_i, Y_i)$ , a candidate split  $A \mid A^c$  is chosen by maximizing an empirical separation score between the child responses. CART uses the between-group mean separation

$$\frac{n_L n_R}{n^2} \|\bar{Y}_L - \bar{Y}_R\|^2,$$

which is the identity-kernel mean-embedding score. Random Fourier feature MMD splitting replaces  $Y$  by a complex feature map  $\psi(Y)$  and scores

$$\frac{n_L n_R}{n^2} \left\| \frac{1}{n_L} \sum_{i \in L} \psi(Y_i) - \frac{1}{n_R} \sum_{i \in R} \psi(Y_i) \right\|^2.$$

With enough features and enough data this score can detect changes beyond the mean. At finite sample sizes it also spends split-selection capacity on noisy higher-order coordinates.

This suggests the working characterization tested below. CART concentrates on the dominant location signal. MMD spreads sensitivity across location and higher-order distributional structure. MMD should win only when the additional structure is strong enough, or the sample size large enough, for its split-time signal-to-noise ratio to exceed the cost of estimating it.

We compare five criteria: CART; isotropic Gaussian RFF MMD (`mmd_rff`); diagonal-bandwidth MMD (`anisotropic_mmd`); an overcomplete RFF pool with per-split top- $k$  frequency selection (`adaptive_mmd`); and a non-kernel sliced-Wasserstein criterion.<sup>1</sup>

## 3 Evidence

The experiments were run as paired-seed comparisons: the train/test split, forest seeds, tree hyperparameters, honesty setting, and cutpoint grid were held fixed while varying only the split criterion. Metrics are RMSE for the conditional mean and CRPS / energy score for distributional prediction. In every table a negative difference means the row improves over the reference, and the parenthetical is the seed-level win rate.

**Synthetic diagnostics.** The first set of experiments swept the univariate synthetic quantile mechanisms used in the DRF literature (Table 1). When the conditional law carries deliberately strong non-location structure, distributional splitting pays: on `paper_quantile_2` and `paper_quantile_3` with honesty 0.5, every distributional criterion improves CRPS and energy over CART with a 100% win rate while leaving RMSE essentially unchanged. On the more location-driven `paper_quantile_1` the comparison is a wash. Honesty is a visible lever in its own right: the gap between criteria at honesty 0 largely collapses once sample splitting is turned on.

**Real univariate data.** On real univariate benchmarks CART-style splitting is again the stronger default (Table 2). CART wins `wine_quality_white` and `kin8nm` against the MMD family, and `diabetes` ( $n = 300$ ) is too noisy to separate the criteria. The exception worth isolating is California housing, where the MMD family beats CART consistently (isotropic MMD improves CRPS by  $\approx 0.003$  at a 100% win rate, and adaptive MMD roughly doubles that). That case confounds

<sup>1</sup>Diagonal-bandwidth MMD is reported only for multivariate responses. On a one-dimensional response it is not a distinct estimator: the coordinatewise median bandwidth coincides with the Euclidean median, and the length-one frequency draw is the same realization as the isotropic scalar-scale draw, so `anisotropic_mmd` reproduces `mmd_rff` bit-for-bit and is omitted from the scalar tables.

Table 1: Paired differences (criterion – `cart`). Negative means the row improves over the reference; values are mean  $\pm$  paired SE with seed-level win rate.

| dataset          | $n_{tr}$ | honesty | criterion          | pairs | $\Delta RMSE$               | $\Delta CRPS$               | $\Delta energy$             |
|------------------|----------|---------|--------------------|-------|-----------------------------|-----------------------------|-----------------------------|
| paper_quantile_1 | 1500     | 0       | adaptive_mmd       | 10    | -0.0069 $\pm$ 0.0012 (90%)  | -0.0032 $\pm$ 0.0006 (100%) | -0.0032 $\pm$ 0.0006 (100%) |
| paper_quantile_1 | 1500     | 0       | mmd_rff            | 10    | -0.0026 $\pm$ 0.0011 (80%)  | -0.0006 $\pm$ 0.0006 (60%)  | -0.0006 $\pm$ 0.0006 (60%)  |
| paper_quantile_1 | 1500     | 0       | sliced_wasserstein | 10    | -0.0011 $\pm$ 0.0008 (60%)  | -0.0001 $\pm$ 0.0005 (40%)  | -0.0001 $\pm$ 0.0005 (40%)  |
| paper_quantile_1 | 1500     | 0.5     | adaptive_mmd       | 10    | +0.0003 $\pm$ 0.0006 (20%)  | +0.0001 $\pm$ 0.0003 (30%)  | +0.0001 $\pm$ 0.0003 (30%)  |
| paper_quantile_1 | 1500     | 0.5     | mmd_rff            | 10    | +0.0002 $\pm$ 0.0003 (40%)  | +0.0001 $\pm$ 0.0002 (30%)  | +0.0001 $\pm$ 0.0002 (30%)  |
| paper_quantile_1 | 1500     | 0.5     | sliced_wasserstein | 10    | -0.0000 $\pm$ 0.0004 (50%)  | -0.0001 $\pm$ 0.0002 (50%)  | -0.0001 $\pm$ 0.0002 (50%)  |
| paper_quantile_2 | 1500     | 0       | adaptive_mmd       | 10    | -0.0082 $\pm$ 0.0018 (100%) | -0.0089 $\pm$ 0.0011 (100%) | -0.0089 $\pm$ 0.0011 (100%) |
| paper_quantile_2 | 1500     | 0       | mmd_rff            | 10    | -0.0026 $\pm$ 0.0024 (70%)  | -0.0060 $\pm$ 0.0014 (90%)  | -0.0060 $\pm$ 0.0014 (90%)  |
| paper_quantile_2 | 1500     | 0       | sliced_wasserstein | 10    | +0.0010 $\pm$ 0.0016 (30%)  | -0.0054 $\pm$ 0.0010 (100%) | -0.0054 $\pm$ 0.0010 (100%) |
| paper_quantile_2 | 1500     | 0.5     | adaptive_mmd       | 10    | +0.0001 $\pm$ 0.0005 (70%)  | -0.0105 $\pm$ 0.0011 (100%) | -0.0105 $\pm$ 0.0011 (100%) |
| paper_quantile_2 | 1500     | 0.5     | mmd_rff            | 10    | +0.0003 $\pm$ 0.0004 (40%)  | -0.0105 $\pm$ 0.0009 (100%) | -0.0105 $\pm$ 0.0009 (100%) |
| paper_quantile_2 | 1500     | 0.5     | sliced_wasserstein | 10    | +0.0007 $\pm$ 0.0006 (40%)  | -0.0105 $\pm$ 0.0011 (100%) | -0.0105 $\pm$ 0.0011 (100%) |
| paper_quantile_3 | 1500     | 0       | adaptive_mmd       | 10    | -0.0057 $\pm$ 0.0017 (90%)  | -0.0043 $\pm$ 0.0015 (80%)  | -0.0043 $\pm$ 0.0015 (80%)  |
| paper_quantile_3 | 1500     | 0       | mmd_rff            | 10    | -0.0018 $\pm$ 0.0015 (80%)  | -0.0023 $\pm$ 0.0012 (80%)  | -0.0023 $\pm$ 0.0012 (80%)  |
| paper_quantile_3 | 1500     | 0       | sliced_wasserstein | 10    | +0.0015 $\pm$ 0.0015 (40%)  | -0.0021 $\pm$ 0.0009 (70%)  | -0.0021 $\pm$ 0.0009 (70%)  |
| paper_quantile_3 | 1500     | 0.5     | adaptive_mmd       | 10    | +0.0007 $\pm$ 0.0005 (30%)  | -0.0034 $\pm$ 0.0008 (100%) | -0.0034 $\pm$ 0.0008 (100%) |
| paper_quantile_3 | 1500     | 0.5     | mmd_rff            | 10    | +0.0003 $\pm$ 0.0004 (50%)  | -0.0033 $\pm$ 0.0006 (100%) | -0.0033 $\pm$ 0.0006 (100%) |
| paper_quantile_3 | 1500     | 0.5     | sliced_wasserstein | 10    | +0.0004 $\pm$ 0.0003 (30%)  | -0.0030 $\pm$ 0.0004 (100%) | -0.0030 $\pm$ 0.0004 (100%) |

two explanations: the dataset is large, and housing prices may contain estimable heteroskedastic structure. A subsample experiment over  $n \in \{2000, 4000, 8000, 16000\}$  separates these readings (Figure 1). The paired MMD–CART advantage is already present at small  $n$  and is sustained as  $n$  grows, which favors the strong-distributional-structure reading over a pure estimability threshold.

Table 2: Paired differences (criterion – `cart`). Negative means the row improves over the reference; values are mean  $\pm$  paired SE with seed-level win rate.

| dataset            | $n_{tr}$ | honesty | criterion          | pairs | $\Delta RMSE$               | $\Delta CRPS$               | $\Delta energy$             |
|--------------------|----------|---------|--------------------|-------|-----------------------------|-----------------------------|-----------------------------|
| california_housing | 3000     | 0       | adaptive_mmd       | 10    | -0.0051 $\pm$ 0.0019 (80%)  | -0.0053 $\pm$ 0.0008 (100%) | -0.0053 $\pm$ 0.0008 (100%) |
| california_housing | 3000     | 0       | mmd_rff            | 10    | -0.0032 $\pm$ 0.0006 (90%)  | -0.0026 $\pm$ 0.0003 (100%) | -0.0026 $\pm$ 0.0003 (100%) |
| california_housing | 3000     | 0       | sliced_wasserstein | 10    | -0.0005 $\pm$ 0.0008 (70%)  | -0.0008 $\pm$ 0.0004 (80%)  | -0.0008 $\pm$ 0.0004 (80%)  |
| california_housing | 3000     | 0.5     | adaptive_mmd       | 10    | -0.0075 $\pm$ 0.0014 (90%)  | -0.0055 $\pm$ 0.0007 (100%) | -0.0055 $\pm$ 0.0007 (100%) |
| california_housing | 3000     | 0.5     | mmd_rff            | 10    | -0.0045 $\pm$ 0.0006 (100%) | -0.0029 $\pm$ 0.0003 (100%) | -0.0029 $\pm$ 0.0003 (100%) |
| california_housing | 3000     | 0.5     | sliced_wasserstein | 10    | -0.0017 $\pm$ 0.0006 (80%)  | -0.0010 $\pm$ 0.0003 (80%)  | -0.0010 $\pm$ 0.0003 (80%)  |
| diabetes           | 300      | 0       | adaptive_mmd       | 10    | +0.1814 $\pm$ 0.1581 (50%)  | +0.2474 $\pm$ 0.0867 (20%)  | +0.2474 $\pm$ 0.0867 (20%)  |
| diabetes           | 300      | 0       | mmd_rff            | 10    | +0.0565 $\pm$ 0.0947 (50%)  | +0.1077 $\pm$ 0.0572 (30%)  | +0.1077 $\pm$ 0.0572 (30%)  |
| diabetes           | 300      | 0       | sliced_wasserstein | 10    | +0.0831 $\pm$ 0.1027 (50%)  | +0.0811 $\pm$ 0.0684 (30%)  | +0.0811 $\pm$ 0.0684 (30%)  |
| diabetes           | 300      | 0.5     | adaptive_mmd       | 10    | +0.1425 $\pm$ 0.0994 (30%)  | +0.1065 $\pm$ 0.0549 (30%)  | +0.1065 $\pm$ 0.0549 (30%)  |
| diabetes           | 300      | 0.5     | mmd_rff            | 10    | +0.0222 $\pm$ 0.0595 (50%)  | +0.0176 $\pm$ 0.0363 (50%)  | +0.0176 $\pm$ 0.0363 (50%)  |
| diabetes           | 300      | 0.5     | sliced_wasserstein | 10    | +0.0401 $\pm$ 0.0655 (40%)  | +0.0220 $\pm$ 0.0368 (50%)  | +0.0220 $\pm$ 0.0368 (50%)  |
| kin8nm             | 3000     | 0       | adaptive_mmd       | 10    | +0.0069 $\pm$ 0.0003 (0%)   | +0.0031 $\pm$ 0.0001 (0%)   | +0.0031 $\pm$ 0.0001 (0%)   |
| kin8nm             | 3000     | 0       | mmd_rff            | 10    | +0.0014 $\pm$ 0.0002 (10%)  | +0.0005 $\pm$ 0.0001 (10%)  | +0.0005 $\pm$ 0.0001 (10%)  |
| kin8nm             | 3000     | 0       | sliced_wasserstein | 10    | -0.0016 $\pm$ 0.0002 (100%) | -0.0008 $\pm$ 0.0001 (100%) | -0.0008 $\pm$ 0.0001 (100%) |
| kin8nm             | 3000     | 0.5     | adaptive_mmd       | 10    | +0.0035 $\pm$ 0.0003 (0%)   | +0.0014 $\pm$ 0.0001 (0%)   | +0.0014 $\pm$ 0.0001 (0%)   |
| kin8nm             | 3000     | 0.5     | mmd_rff            | 10    | +0.0007 $\pm$ 0.0001 (0%)   | +0.0001 $\pm$ 0.0001 (40%)  | +0.0001 $\pm$ 0.0001 (40%)  |
| kin8nm             | 3000     | 0.5     | sliced_wasserstein | 10    | -0.0006 $\pm$ 0.0001 (100%) | -0.0004 $\pm$ 0.0000 (100%) | -0.0004 $\pm$ 0.0000 (100%) |
| wine_quality_white | 3000     | 0       | adaptive_mmd       | 10    | +0.0191 $\pm$ 0.0015 (0%)   | +0.0075 $\pm$ 0.0009 (0%)   | +0.0075 $\pm$ 0.0009 (0%)   |
| wine_quality_white | 3000     | 0       | mmd_rff            | 10    | +0.0068 $\pm$ 0.0009 (0%)   | +0.0021 $\pm$ 0.0005 (10%)  | +0.0021 $\pm$ 0.0005 (10%)  |
| wine_quality_white | 3000     | 0       | sliced_wasserstein | 10    | +0.0072 $\pm$ 0.0010 (0%)   | +0.0024 $\pm$ 0.0005 (10%)  | +0.0024 $\pm$ 0.0005 (10%)  |
| wine_quality_white | 3000     | 0.5     | adaptive_mmd       | 10    | +0.0073 $\pm$ 0.0009 (0%)   | +0.0035 $\pm$ 0.0005 (0%)   | +0.0035 $\pm$ 0.0005 (0%)   |
| wine_quality_white | 3000     | 0.5     | mmd_rff            | 10    | +0.0026 $\pm$ 0.0005 (0%)   | +0.0011 $\pm$ 0.0002 (10%)  | +0.0011 $\pm$ 0.0002 (10%)  |
| wine_quality_white | 3000     | 0.5     | sliced_wasserstein | 10    | +0.0043 $\pm$ 0.0006 (0%)   | +0.0023 $\pm$ 0.0003 (0%)   | +0.0023 $\pm$ 0.0003 (0%)   |

**Do the MMD refinements help?** Table 3 references plain isotropic MMD rather than CART and asks whether any refinement is a better distributional baseline. None is. Adaptive frequency selection helps on California but hurts on `kin8nm` and `wine`; sliced Wasserstein helps on `kin8nm` but loses elsewhere; on the multivariate cells the refinements scatter around zero relative to `mmd_rff` with




Figure 1: California housing subsample curve: paired CRPS difference (criterion – CART) versus  $n_{\text{train}}$ . Points below zero mean the criterion beats CART. The MMD family is below CART across the whole range rather than only at large  $n$ .

no consistent sign. This is the evidence that ordinary RFF-MMD is already the right distributional default: the extra machinery does not move the frontier. Adaptive-frequency MMD in particular was included as a rescue attempt—if uniform RFF averaging were hiding useful signal, per-split selection should expose it—and it does not change the conclusion, which points away from a simple tuning failure and toward the finite-sample allocation explanation.

**Multivariate responses.** The regime that most naturally motivates distributional forests is multivariate response, and here distributional splitting clearly earns its keep (Table 4). On the building-energy real benchmark (`enb`,  $d = 2$ ) and on two synthetic multivariate mechanisms, the MMD family beats CART on the energy score with high win rates. The cleanest case is `paper_copula`, a pure-dependence shift with little marginal mean signal: every distributional criterion improves the energy score over CART by a wide margin at a 100% win rate, while CRPS and RMSE—which see the marginals only—barely move. This is exactly the structure CART cannot see and a distributional criterion can.

## 4 Characterization

The empirical pattern can be summarized as follows.

Distributional splitting helps when non-location conditional structure is present and estimable at node scale. For scalar responses in mean-dominated finite samples, CART-style splitting is the better default; among distributional criteria, ordinary isotropic MMD is already close to best in class, and the anisotropic, adaptive, and sliced-Wasserstein refinements do not systematically improve on it. For multivariate responses, where dependence structure is invisible to mean splitting, distributional criteria are the right tool.

This statement is intentionally scoped. It is not a claim that distributional forests are inferior to quantile regression forests, nor that MMD is a poor criterion in general. The core motivation for

Table 3: Paired differences (criterion  $-$  mmd\_rff). Negative means the row improves over the reference; values are mean  $\pm$  paired SE with seed-level win rate.

| dataset                        | $n_{tr}$ | honesty | criterion          | pairs | $\Delta$ RMSE               | $\Delta$ CRPS               | $\Delta$ energy             |
|--------------------------------|----------|---------|--------------------|-------|-----------------------------|-----------------------------|-----------------------------|
| california_housing             | 3000     | 0       | adaptive_mmd       | 10    | -0.0020 $\pm$ 0.0015 (60%)  | -0.0028 $\pm$ 0.0007 (100%) | -0.0028 $\pm$ 0.0007 (100%) |
| california_housing             | 3000     | 0       | cart               | 10    | +0.0032 $\pm$ 0.0006 (10%)  | +0.0026 $\pm$ 0.0003 (0%)   | +0.0026 $\pm$ 0.0003 (0%)   |
| california_housing             | 3000     | 0       | sliced_wasserstein | 10    | +0.0026 $\pm$ 0.0011 (30%)  | +0.0017 $\pm$ 0.0005 (10%)  | +0.0017 $\pm$ 0.0005 (10%)  |
| california_housing             | 3000     | 0.5     | adaptive_mmd       | 10    | -0.0031 $\pm$ 0.0011 (80%)  | -0.0026 $\pm$ 0.0006 (90%)  | -0.0026 $\pm$ 0.0006 (90%)  |
| california_housing             | 3000     | 0.5     | cart               | 10    | +0.0045 $\pm$ 0.0006 (0%)   | +0.0029 $\pm$ 0.0003 (0%)   | +0.0029 $\pm$ 0.0003 (0%)   |
| california_housing             | 3000     | 0.5     | sliced_wasserstein | 10    | +0.0027 $\pm$ 0.0010 (30%)  | +0.0019 $\pm$ 0.0005 (20%)  | +0.0019 $\pm$ 0.0005 (20%)  |
| diabetes                       | 300      | 0       | adaptive_mmd       | 10    | +0.1249 $\pm$ 0.1012 (40%)  | +0.1397 $\pm$ 0.0634 (30%)  | +0.1397 $\pm$ 0.0634 (30%)  |
| diabetes                       | 300      | 0       | cart               | 10    | -0.0565 $\pm$ 0.0947 (50%)  | -0.1077 $\pm$ 0.0572 (70%)  | -0.1077 $\pm$ 0.0572 (70%)  |
| diabetes                       | 300      | 0       | sliced_wasserstein | 10    | +0.0266 $\pm$ 0.1016 (70%)  | -0.0266 $\pm$ 0.0641 (70%)  | -0.0266 $\pm$ 0.0641 (70%)  |
| diabetes                       | 300      | 0.5     | adaptive_mmd       | 10    | +0.1203 $\pm$ 0.0578 (30%)  | +0.0890 $\pm$ 0.0323 (40%)  | +0.0890 $\pm$ 0.0323 (40%)  |
| diabetes                       | 300      | 0.5     | cart               | 10    | -0.0222 $\pm$ 0.0595 (50%)  | -0.0176 $\pm$ 0.0363 (50%)  | -0.0176 $\pm$ 0.0363 (50%)  |
| diabetes                       | 300      | 0.5     | sliced_wasserstein | 10    | +0.0178 $\pm$ 0.0316 (40%)  | +0.0045 $\pm$ 0.0233 (50%)  | +0.0045 $\pm$ 0.0233 (50%)  |
| enb                            | 576      | 0       | adaptive_mmd       | 5     | -0.0034 $\pm$ 0.0018 (80%)  | -0.0017 $\pm$ 0.0010 (80%)  | -0.0023 $\pm$ 0.0014 (80%)  |
| enb                            | 576      | 0       | anisotropic_mmd    | 5     | -0.0021 $\pm$ 0.0025 (60%)  | -0.0010 $\pm$ 0.0010 (40%)  | -0.0013 $\pm$ 0.0015 (60%)  |
| enb                            | 576      | 0       | cart               | 5     | +0.0027 $\pm$ 0.0027 (40%)  | +0.0017 $\pm$ 0.0011 (20%)  | +0.0024 $\pm$ 0.0017 (20%)  |
| enb                            | 576      | 0       | sliced_wasserstein | 5     | +0.0090 $\pm$ 0.0032 (0%)   | +0.0029 $\pm$ 0.0013 (20%)  | +0.0040 $\pm$ 0.0018 (20%)  |
| enb                            | 576      | 0.5     | adaptive_mmd       | 5     | -0.0064 $\pm$ 0.0044 (60%)  | -0.0025 $\pm$ 0.0016 (80%)  | -0.0036 $\pm$ 0.0024 (100%) |
| enb                            | 576      | 0.5     | anisotropic_mmd    | 5     | -0.0007 $\pm$ 0.0016 (60%)  | -0.0001 $\pm$ 0.0004 (40%)  | -0.0004 $\pm$ 0.0006 (60%)  |
| enb                            | 576      | 0.5     | cart               | 5     | +0.0071 $\pm$ 0.0037 (20%)  | +0.0027 $\pm$ 0.0013 (0%)   | +0.0036 $\pm$ 0.0017 (0%)   |
| enb                            | 576      | 0.5     | sliced_wasserstein | 5     | +0.0089 $\pm$ 0.0031 (0%)   | +0.0047 $\pm$ 0.0010 (0%)   | +0.0058 $\pm$ 0.0013 (0%)   |
| kin8nm                         | 3000     | 0       | adaptive_mmd       | 10    | +0.0056 $\pm$ 0.0002 (0%)   | +0.0026 $\pm$ 0.0001 (0%)   | +0.0026 $\pm$ 0.0001 (0%)   |
| kin8nm                         | 3000     | 0       | cart               | 10    | -0.0014 $\pm$ 0.0002 (90%)  | -0.0005 $\pm$ 0.0001 (90%)  | -0.0005 $\pm$ 0.0001 (90%)  |
| kin8nm                         | 3000     | 0       | sliced_wasserstein | 10    | -0.0030 $\pm$ 0.0003 (100%) | -0.0013 $\pm$ 0.0001 (100%) | -0.0013 $\pm$ 0.0001 (100%) |
| kin8nm                         | 3000     | 0.5     | adaptive_mmd       | 10    | +0.0028 $\pm$ 0.0002 (0%)   | +0.0013 $\pm$ 0.0001 (0%)   | +0.0013 $\pm$ 0.0001 (0%)   |
| kin8nm                         | 3000     | 0.5     | cart               | 10    | -0.0007 $\pm$ 0.0001 (100%) | -0.0001 $\pm$ 0.0001 (60%)  | -0.0001 $\pm$ 0.0001 (60%)  |
| kin8nm                         | 3000     | 0.5     | sliced_wasserstein | 10    | -0.0013 $\pm$ 0.0001 (100%) | -0.0005 $\pm$ 0.0001 (100%) | -0.0005 $\pm$ 0.0001 (100%) |
| paper_copula                   | 3750     | 0       | adaptive_mmd       | 5     | -0.0007 $\pm$ 0.0003 (80%)  | -0.0003 $\pm$ 0.0002 (60%)  | -0.0010 $\pm$ 0.0005 (100%) |
| paper_copula                   | 3750     | 0       | anisotropic_mmd    | 5     | -0.0008 $\pm$ 0.0006 (60%)  | +0.0002 $\pm$ 0.0004 (40%)  | -0.0005 $\pm$ 0.0008 (60%)  |
| paper_copula                   | 3750     | 0       | cart               | 5     | +0.0009 $\pm$ 0.0005 (40%)  | +0.0004 $\pm$ 0.0003 (40%)  | +0.0045 $\pm$ 0.0007 (0%)   |
| paper_copula                   | 3750     | 0       | sliced_wasserstein | 5     | +0.0003 $\pm$ 0.0007 (40%)  | +0.0002 $\pm$ 0.0004 (60%)  | -0.0002 $\pm$ 0.0009 (60%)  |
| paper_copula                   | 3750     | 0.5     | adaptive_mmd       | 5     | -0.0001 $\pm$ 0.0002 (40%)  | -0.0000 $\pm$ 0.0001 (40%)  | -0.0002 $\pm$ 0.0003 (60%)  |
| paper_copula                   | 3750     | 0.5     | anisotropic_mmd    | 5     | +0.0003 $\pm$ 0.0003 (20%)  | +0.0003 $\pm$ 0.0001 (20%)  | -0.0005 $\pm$ 0.0003 (80%)  |
| paper_copula                   | 3750     | 0.5     | cart               | 5     | -0.0002 $\pm$ 0.0003 (60%)  | -0.0001 $\pm$ 0.0001 (60%)  | +0.0038 $\pm$ 0.0003 (0%)   |
| paper_copula                   | 3750     | 0.5     | sliced_wasserstein | 5     | -0.0002 $\pm$ 0.0003 (60%)  | -0.0000 $\pm$ 0.0002 (60%)  | -0.0008 $\pm$ 0.0004 (80%)  |
| paper_heterogeneous_regression | 3750     | 0       | adaptive_mmd       | 5     | -0.0008 $\pm$ 0.0010 (80%)  | +0.0001 $\pm$ 0.0008 (60%)  | +0.0001 $\pm$ 0.0011 (60%)  |
| paper_heterogeneous_regression | 3750     | 0       | anisotropic_mmd    | 5     | -0.0005 $\pm$ 0.0007 (60%)  | +0.0005 $\pm$ 0.0005 (40%)  | +0.0008 $\pm$ 0.0007 (20%)  |
| paper_heterogeneous_regression | 3750     | 0       | cart               | 5     | -0.0012 $\pm$ 0.0018 (60%)  | +0.0006 $\pm$ 0.0008 (20%)  | +0.0010 $\pm$ 0.0012 (20%)  |
| paper_heterogeneous_regression | 3750     | 0       | sliced_wasserstein | 5     | -0.0012 $\pm$ 0.0007 (80%)  | -0.0005 $\pm$ 0.0003 (80%)  | -0.0008 $\pm$ 0.0005 (80%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | adaptive_mmd       | 5     | +0.0020 $\pm$ 0.0012 (20%)  | +0.0004 $\pm$ 0.0006 (20%)  | +0.0007 $\pm$ 0.0010 (20%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | anisotropic_mmd    | 5     | +0.0007 $\pm$ 0.0008 (40%)  | +0.0000 $\pm$ 0.0003 (40%)  | +0.0002 $\pm$ 0.0006 (40%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | cart               | 5     | +0.0023 $\pm$ 0.0010 (0%)   | +0.0042 $\pm$ 0.0004 (0%)   | +0.0069 $\pm$ 0.0008 (0%)   |
| paper_heterogeneous_regression | 3750     | 0.5     | sliced_wasserstein | 5     | +0.0013 $\pm$ 0.0005 (20%)  | +0.0012 $\pm$ 0.0002 (0%)   | +0.0019 $\pm$ 0.0003 (0%)   |
| wine_quality_white             | 3000     | 0       | adaptive_mmd       | 10    | +0.0123 $\pm$ 0.0006 (0%)   | +0.0054 $\pm$ 0.0004 (0%)   | +0.0054 $\pm$ 0.0004 (0%)   |
| wine_quality_white             | 3000     | 0       | cart               | 10    | -0.0068 $\pm$ 0.0009 (100%) | -0.0021 $\pm$ 0.0005 (90%)  | -0.0021 $\pm$ 0.0005 (90%)  |
| wine_quality_white             | 3000     | 0       | sliced_wasserstein | 10    | +0.0004 $\pm$ 0.0013 (50%)  | +0.0003 $\pm$ 0.0007 (60%)  | +0.0003 $\pm$ 0.0007 (60%)  |
| wine_quality_white             | 3000     | 0.5     | adaptive_mmd       | 10    | +0.0047 $\pm$ 0.0005 (0%)   | +0.0024 $\pm$ 0.0003 (0%)   | +0.0024 $\pm$ 0.0003 (0%)   |
| wine_quality_white             | 3000     | 0.5     | cart               | 10    | -0.0026 $\pm$ 0.0005 (100%) | -0.0011 $\pm$ 0.0002 (90%)  | -0.0011 $\pm$ 0.0002 (90%)  |
| wine_quality_white             | 3000     | 0.5     | sliced_wasserstein | 10    | +0.0017 $\pm$ 0.0006 (20%)  | +0.0012 $\pm$ 0.0004 (20%)  | +0.0012 $\pm$ 0.0004 (20%)  |

DRF is multivariate response distributions, where CART/QRF-style univariate orderings are not a complete solution, and the multivariate evidence above supports that motivation directly. The negative result is about the common univariate tabular setting: when the target is scalar and the conditional law is mostly a shifted version of itself, a distributional split criterion is often paying variance to chase weak structure.

## 5 Implications

For practice, the recommendation is simple. Use CART-style splitting as the univariate default, then reach for a distributional criterion—ordinary isotropic MMD, without further refinement—when the response is multivariate or the scalar data show clear heteroskedasticity, tail behavior, or non-location structure. When distributional splitting is used, paired-seed comparisons and honesty settings should be reported: in these experiments, honesty can be as important as the criterion itself.

Table 4: Paired differences (criterion  $-$  `cart`). Negative means the row improves over the reference; values are mean  $\pm$  paired SE with seed-level win rate.

| dataset                        | $n_{tr}$ | honesty | criterion          | pairs | $\Delta_{energy}$           | $\Delta_{CRPS}$             | $\Delta_{RMSE}$             |
|--------------------------------|----------|---------|--------------------|-------|-----------------------------|-----------------------------|-----------------------------|
| enb                            | 576      | 0       | adaptive_mmd       | 5     | -0.0047 $\pm$ 0.0026 (100%) | -0.0033 $\pm$ 0.0018 (100%) | -0.0062 $\pm$ 0.0038 (80%)  |
| enb                            | 576      | 0       | anisotropic_mmd    | 5     | -0.0037 $\pm$ 0.0024 (80%)  | -0.0026 $\pm$ 0.0017 (80%)  | -0.0048 $\pm$ 0.0041 (80%)  |
| enb                            | 576      | 0       | mmd_rff            | 5     | -0.0024 $\pm$ 0.0017 (80%)  | -0.0017 $\pm$ 0.0011 (80%)  | -0.0027 $\pm$ 0.0027 (60%)  |
| enb                            | 576      | 0       | sliced_wasserstein | 5     | +0.0015 $\pm$ 0.0022 (40%)  | +0.0012 $\pm$ 0.0016 (40%)  | +0.0063 $\pm$ 0.0040 (20%)  |
| enb                            | 576      | 0.5     | adaptive_mmd       | 5     | -0.0072 $\pm$ 0.0031 (100%) | -0.0052 $\pm$ 0.0022 (100%) | -0.0135 $\pm$ 0.0067 (100%) |
| enb                            | 576      | 0.5     | anisotropic_mmd    | 5     | -0.0041 $\pm$ 0.0019 (80%)  | -0.0028 $\pm$ 0.0014 (80%)  | -0.0078 $\pm$ 0.0047 (80%)  |
| enb                            | 576      | 0.5     | mmd_rff            | 5     | -0.0036 $\pm$ 0.0017 (100%) | -0.0027 $\pm$ 0.0013 (100%) | -0.0071 $\pm$ 0.0037 (80%)  |
| enb                            | 576      | 0.5     | sliced_wasserstein | 5     | +0.0022 $\pm$ 0.0027 (40%)  | +0.0020 $\pm$ 0.0020 (20%)  | +0.0018 $\pm$ 0.0058 (60%)  |
| paper_copula                   | 3750     | 0       | adaptive_mmd       | 5     | -0.0054 $\pm$ 0.0010 (100%) | -0.0007 $\pm$ 0.0004 (80%)  | -0.0016 $\pm$ 0.0008 (100%) |
| paper_copula                   | 3750     | 0       | anisotropic_mmd    | 5     | -0.0050 $\pm$ 0.0014 (100%) | -0.0002 $\pm$ 0.0006 (60%)  | -0.0018 $\pm$ 0.0011 (60%)  |
| paper_copula                   | 3750     | 0       | mmd_rff            | 5     | -0.0045 $\pm$ 0.0007 (100%) | -0.0004 $\pm$ 0.0003 (60%)  | -0.0009 $\pm$ 0.0005 (60%)  |
| paper_copula                   | 3750     | 0       | sliced_wasserstein | 5     | -0.0047 $\pm$ 0.0010 (100%) | -0.0002 $\pm$ 0.0004 (80%)  | -0.0006 $\pm$ 0.0006 (80%)  |
| paper_copula                   | 3750     | 0.5     | adaptive_mmd       | 5     | -0.0040 $\pm$ 0.0004 (100%) | +0.0001 $\pm$ 0.0001 (20%)  | +0.0001 $\pm$ 0.0003 (60%)  |
| paper_copula                   | 3750     | 0.5     | anisotropic_mmd    | 5     | -0.0043 $\pm$ 0.0002 (100%) | +0.0004 $\pm$ 0.0002 (20%)  | +0.0004 $\pm$ 0.0003 (20%)  |
| paper_copula                   | 3750     | 0.5     | mmd_rff            | 5     | -0.0038 $\pm$ 0.0003 (100%) | +0.0001 $\pm$ 0.0001 (40%)  | +0.0002 $\pm$ 0.0003 (40%)  |
| paper_copula                   | 3750     | 0.5     | sliced_wasserstein | 5     | -0.0046 $\pm$ 0.0002 (100%) | +0.0001 $\pm$ 0.0001 (20%)  | +0.0000 $\pm$ 0.0001 (60%)  |
| paper_heterogeneous_regression | 3750     | 0       | adaptive_mmd       | 5     | -0.0009 $\pm$ 0.0012 (80%)  | -0.0006 $\pm$ 0.0008 (80%)  | +0.0004 $\pm$ 0.0022 (40%)  |
| paper_heterogeneous_regression | 3750     | 0       | anisotropic_mmd    | 5     | -0.0002 $\pm$ 0.0010 (40%)  | -0.0001 $\pm$ 0.0006 (40%)  | +0.0008 $\pm$ 0.0020 (40%)  |
| paper_heterogeneous_regression | 3750     | 0       | mmd_rff            | 5     | -0.0010 $\pm$ 0.0012 (80%)  | -0.0006 $\pm$ 0.0008 (80%)  | +0.0012 $\pm$ 0.0018 (40%)  |
| paper_heterogeneous_regression | 3750     | 0       | sliced_wasserstein | 5     | -0.0017 $\pm$ 0.0008 (80%)  | -0.0011 $\pm$ 0.0005 (80%)  | -0.0000 $\pm$ 0.0015 (40%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | adaptive_mmd       | 5     | -0.0062 $\pm$ 0.0012 (100%) | -0.0039 $\pm$ 0.0007 (100%) | -0.0003 $\pm$ 0.0018 (20%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | anisotropic_mmd    | 5     | -0.0067 $\pm$ 0.0009 (100%) | -0.0042 $\pm$ 0.0005 (100%) | -0.0016 $\pm$ 0.0014 (80%)  |
| paper_heterogeneous_regression | 3750     | 0.5     | mmd_rff            | 5     | -0.0069 $\pm$ 0.0008 (100%) | -0.0042 $\pm$ 0.0004 (100%) | -0.0023 $\pm$ 0.0010 (100%) |
| paper_heterogeneous_regression | 3750     | 0.5     | sliced_wasserstein | 5     | -0.0050 $\pm$ 0.0007 (100%) | -0.0031 $\pm$ 0.0004 (100%) | -0.0010 $\pm$ 0.0007 (80%)  |

All of the above is implemented in the open-source `drforest` library, which provides the honest forest, every criterion compared here, and the post-hoc shrinkage variants behind one interface; its Rust-backed split search ran the full paired sweep in this note at a cost that makes such characterizations routine rather than a project in themselves.

For research, the next useful theory is not a generic dominance theorem. The useful object is a finite-sample comparison of split-score signal-to-noise: mean separation versus kernel mean separation under local alternatives. Such a calculation would turn the empirical characterization above into a sharp threshold statement.

## 6 Limitations

The scalar experiments compare CART-style splitting to MMD splitting inside the same forest implementation; a full external QRF package comparison is a separate baseline. The multivariate evidence, while decisive on the energy score, uses a small number of mechanisms and one real benchmark, and a broader multivariate suite would strengthen the positive side of the characterization.

## A Shrinkage Frontier

Post-hoc shrinkage of forest weights is a separate lever from the split criterion. Table 5 pairs each shrinkage variant against raw weights within a fixed criterion and forest, so the difference isolates the shrinkage effect. Shrinkage does not change the criterion ranking: marginal shrinkage is at best a small CRPS improvement on one synthetic mechanism, parent shrinkage runs to large intensity and consistently hurts, and the pattern is the same across CART, MMD, and sliced-Wasserstein forests. Shrinkage is not the missing ingredient that would let the MMD variants dominate CART on scalar data.

Table 5: Shrinkage frontier: paired differences (variant – raw) within a fixed criterion and forest. Negative means shrinkage improves over raw weights;  $\bar{\alpha}$  is the mean shrinkage intensity.

| dataset          | criterion          | honesty | variant        | pairs | $\bar{\alpha}$ | $\Delta$ RMSE               | $\Delta$ CRPS              | $\Delta$ energy            |
|------------------|--------------------|---------|----------------|-------|----------------|-----------------------------|----------------------------|----------------------------|
| enb              | cart               | 0.5     | marginal_kmse  | 5     | 0.003          | -0.0001 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0012 $\pm$ 0.0002 (0%)  |
| enb              | cart               | 0.5     | marginal_stein | 5     | 0.003          | -0.0001 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0012 $\pm$ 0.0002 (0%)  |
| enb              | cart               | 0.5     | parent_kmse    | 5     | 0.282          | +0.0472 $\pm$ 0.0053 (0%)   | +0.0351 $\pm$ 0.0017 (0%)  | +0.0509 $\pm$ 0.0022 (0%)  |
| enb              | cart               | 0.5     | parent_stein   | 5     | 0.402          | +0.0612 $\pm$ 0.0084 (0%)   | +0.0494 $\pm$ 0.0029 (0%)  | +0.0719 $\pm$ 0.0039 (0%)  |
| enb              | mmd_rff            | 0.5     | marginal_kmse  | 5     | 0.003          | -0.0001 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0013 $\pm$ 0.0002 (0%)  |
| enb              | mmd_rff            | 0.5     | marginal_stein | 5     | 0.003          | -0.0001 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0013 $\pm$ 0.0002 (0%)  |
| enb              | mmd_rff            | 0.5     | parent_kmse    | 5     | 0.286          | +0.0487 $\pm$ 0.0044 (0%)   | +0.0361 $\pm$ 0.0013 (0%)  | +0.0522 $\pm$ 0.0016 (0%)  |
| enb              | mmd_rff            | 0.5     | parent_stein   | 5     | 0.408          | +0.0623 $\pm$ 0.0070 (0%)   | +0.0507 $\pm$ 0.0023 (0%)  | +0.0735 $\pm$ 0.0030 (0%)  |
| enb              | sliced_wasserstein | 0.5     | marginal_kmse  | 5     | 0.003          | -0.0001 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0012 $\pm$ 0.0002 (0%)  |
| enb              | sliced_wasserstein | 0.5     | marginal_stein | 5     | 0.003          | -0.0002 $\pm$ 0.0008 (40%)  | +0.0009 $\pm$ 0.0001 (0%)  | +0.0012 $\pm$ 0.0002 (0%)  |
| enb              | sliced_wasserstein | 0.5     | parent_kmse    | 5     | 0.275          | +0.0495 $\pm$ 0.0052 (0%)   | +0.0360 $\pm$ 0.0015 (0%)  | +0.0520 $\pm$ 0.0019 (0%)  |
| enb              | sliced_wasserstein | 0.5     | parent_stein   | 5     | 0.396          | +0.0644 $\pm$ 0.0081 (0%)   | +0.0510 $\pm$ 0.0027 (0%)  | +0.0740 $\pm$ 0.0035 (0%)  |
| paper_quantile_2 | cart               | 0.5     | marginal_kmse  | 5     | 0.316          | -0.0007 $\pm$ 0.0002 (100%) | +0.0018 $\pm$ 0.0003 (0%)  | +0.0018 $\pm$ 0.0003 (0%)  |
| paper_quantile_2 | cart               | 0.5     | marginal_stein | 5     | 0.460          | -0.0010 $\pm$ 0.0003 (100%) | +0.0026 $\pm$ 0.0005 (0%)  | +0.0026 $\pm$ 0.0005 (0%)  |
| paper_quantile_2 | cart               | 0.5     | parent_kmse    | 5     | 0.839          | -0.0004 $\pm$ 0.0003 (80%)  | +0.0009 $\pm$ 0.0002 (0%)  | +0.0009 $\pm$ 0.0002 (0%)  |
| paper_quantile_2 | cart               | 0.5     | parent_stein   | 5     | 1.000          | -0.0004 $\pm$ 0.0004 (80%)  | +0.0013 $\pm$ 0.0003 (0%)  | +0.0013 $\pm$ 0.0003 (0%)  |
| paper_quantile_2 | mmd_rff            | 0.5     | marginal_kmse  | 5     | 0.059          | -0.0001 $\pm$ 0.0001 (80%)  | +0.0002 $\pm$ 0.0001 (20%) | +0.0002 $\pm$ 0.0001 (20%) |
| paper_quantile_2 | mmd_rff            | 0.5     | marginal_stein | 5     | 0.066          | -0.0002 $\pm$ 0.0001 (80%)  | +0.0002 $\pm$ 0.0002 (40%) | +0.0002 $\pm$ 0.0002 (40%) |
| paper_quantile_2 | mmd_rff            | 0.5     | parent_kmse    | 5     | 0.903          | -0.0002 $\pm$ 0.0005 (60%)  | -0.0001 $\pm$ 0.0003 (60%) | -0.0001 $\pm$ 0.0003 (60%) |
| paper_quantile_2 | mmd_rff            | 0.5     | parent_stein   | 5     | 1.000          | -0.0003 $\pm$ 0.0005 (60%)  | -0.0002 $\pm$ 0.0003 (60%) | -0.0002 $\pm$ 0.0003 (60%) |
| paper_quantile_2 | sliced_wasserstein | 0.5     | marginal_kmse  | 5     | 0.057          | -0.0002 $\pm$ 0.0001 (80%)  | +0.0002 $\pm$ 0.0001 (20%) | +0.0002 $\pm$ 0.0001 (20%) |
| paper_quantile_2 | sliced_wasserstein | 0.5     | marginal_stein | 5     | 0.064          | -0.0001 $\pm$ 0.0000 (80%)  | +0.0002 $\pm$ 0.0002 (20%) | +0.0002 $\pm$ 0.0002 (20%) |
| paper_quantile_2 | sliced_wasserstein | 0.5     | parent_kmse    | 5     | 0.904          | -0.0000 $\pm$ 0.0003 (60%)  | +0.0001 $\pm$ 0.0002 (40%) | +0.0001 $\pm$ 0.0002 (40%) |
| paper_quantile_2 | sliced_wasserstein | 0.5     | parent_stein   | 5     | 1.000          | -0.0000 $\pm$ 0.0004 (60%)  | +0.0001 $\pm$ 0.0003 (40%) | +0.0001 $\pm$ 0.0003 (40%) |
| shrinkage_toy    | cart               | 0.5     | marginal_kmse  | 5     | 0.040          | +0.0076 $\pm$ 0.0018 (0%)   | +0.0038 $\pm$ 0.0007 (0%)  | +0.0061 $\pm$ 0.0011 (0%)  |
| shrinkage_toy    | cart               | 0.5     | marginal_stein | 5     | 0.042          | +0.0081 $\pm$ 0.0020 (0%)   | +0.0040 $\pm$ 0.0007 (0%)  | +0.0064 $\pm$ 0.0012 (0%)  |
| shrinkage_toy    | cart               | 0.5     | parent_kmse    | 5     | 0.379          | +0.0137 $\pm$ 0.0036 (0%)   | +0.0072 $\pm$ 0.0014 (0%)  | +0.0135 $\pm$ 0.0023 (0%)  |
| shrinkage_toy    | cart               | 0.5     | parent_stein   | 5     | 0.586          | +0.0229 $\pm$ 0.0053 (0%)   | +0.0119 $\pm$ 0.0023 (0%)  | +0.0223 $\pm$ 0.0036 (0%)  |
| shrinkage_toy    | mmd_rff            | 0.5     | marginal_kmse  | 5     | 0.039          | +0.0075 $\pm$ 0.0019 (0%)   | +0.0037 $\pm$ 0.0007 (0%)  | +0.0060 $\pm$ 0.0011 (0%)  |
| shrinkage_toy    | mmd_rff            | 0.5     | marginal_stein | 5     | 0.041          | +0.0079 $\pm$ 0.0020 (0%)   | +0.0039 $\pm$ 0.0007 (0%)  | +0.0063 $\pm$ 0.0012 (0%)  |
| shrinkage_toy    | mmd_rff            | 0.5     | parent_kmse    | 5     | 0.376          | +0.0131 $\pm$ 0.0036 (0%)   | +0.0068 $\pm$ 0.0015 (0%)  | +0.0129 $\pm$ 0.0024 (0%)  |
| shrinkage_toy    | mmd_rff            | 0.5     | parent_stein   | 5     | 0.578          | +0.0216 $\pm$ 0.0054 (0%)   | +0.0112 $\pm$ 0.0024 (0%)  | +0.0211 $\pm$ 0.0037 (0%)  |
| shrinkage_toy    | sliced_wasserstein | 0.5     | marginal_kmse  | 5     | 0.039          | +0.0076 $\pm$ 0.0019 (0%)   | +0.0038 $\pm$ 0.0007 (0%)  | +0.0061 $\pm$ 0.0011 (0%)  |
| shrinkage_toy    | sliced_wasserstein | 0.5     | marginal_stein | 5     | 0.042          | +0.0081 $\pm$ 0.0020 (0%)   | +0.0040 $\pm$ 0.0007 (0%)  | +0.0065 $\pm$ 0.0012 (0%)  |
| shrinkage_toy    | sliced_wasserstein | 0.5     | parent_kmse    | 5     | 0.383          | +0.0131 $\pm$ 0.0034 (0%)   | +0.0068 $\pm$ 0.0013 (0%)  | +0.0130 $\pm$ 0.0021 (0%)  |
| shrinkage_toy    | sliced_wasserstein | 0.5     | parent_stein   | 5     | 0.591          | +0.0220 $\pm$ 0.0051 (0%)   | +0.0114 $\pm$ 0.0022 (0%)  | +0.0216 $\pm$ 0.0034 (0%)  |

## References

- [1] Domagoj Čevič, Loris Michel, Jeffrey Näf, Nicolai Meinshausen, and Peter Bühlmann. Distributional random forests: Heterogeneity adjustment and multivariate distributional regression. *Journal of Machine Learning Research*, 23(333):1–79, 2022.
- [2] Nicolai Meinshausen. Quantile regression forests. *Journal of Machine Learning Research*, 7: 983–999, 2006.